



Generate Verbalized Confidence



Approximate True Confidence



Revise the Completion



Construct Teacher Context



Calculate Reverse KL loss



Use prompt instruction to force the student model to output its own confidence score C in $[0,1]$ alongside the response

Leverage multiple rollouts to calculate the actual success probability (C_{acc}), approximating the "true" confidence

Revise the student's raw response by replacing its original confidence C with the calculated C_{acc} , enforcing confidence

Create a privileged context combining the user question, feedback or corrected demo, and the tampered completion

Compare student and teacher logprobs on the same completion to minimize the distribution gap via reverse KL divergence